

MNIST Handwritten Digit Classifier Using JAX

Project Title

MNIST Handwritten Digit Classifier Using JAX

Description

This is a simple neural network built from scratch using **JAX**, trained to recognize handwritten digits from the **MNIST** dataset (0–9).

- Framework: [JAX](#)
 - Dataset: [MNIST](#)
 - Model: 1 hidden layer neural network
 - Purpose: Learn how neural networks work using forward pass, ReLU, softmax, loss, and gradient descent — all coded manually.
-

What is MNIST?

The **MNIST dataset** contains **70,000 grayscale images** of handwritten digits:

- 60,000 training images
 - 10,000 testing images
 - Each image: 28×28 pixels, values range from 0 to 255
-

How Neural Networks Work (with Forward Pass)

Step-by-step:

1. Input Image:

2. Each image is 28×28 , flattened into a **vector of 784 values** (pixels).

3. Normalized between 0 and 1 by dividing by 255.

4. Weights Initialization:

5. W_1 : (784×128) — connects input to hidden layer

6. b_1 : $(128,)$ — bias for hidden layer

7. W_2 : (128×10) — connects hidden to output (10 digits)

8. b_2 : $(10,)$ — bias for output

9. Forward Pass:

```
# Hidden Layer
z1 = X @ W1 + b1
a1 = relu(z1)

# Output Layer
z2 = a1 @ W2 + b2
y_hat = softmax(z2)
```

- `relu(x)`: ReLU (Rectified Linear Unit) turns negative values to 0.
- `softmax(x)`: Converts logits into probabilities across 10 classes.
- **Prediction:**
 - The digit with the **highest probability** is the predicted class.

Loss Function

```
def cross_entropy(preds, labels):
    one_hot = jnp.eye(10)[labels] # one-hot encode
    return -jnp.sum(one_hot * jnp.log(preds + 1e-7))
```

- Cross-entropy compares predicted probabilities to the actual label.
- Lower loss = better prediction.

Training Loop

1. **Loop for 1000 epochs**
2. Each step:
3. Pick a sample (`X_train[i]`)
4. Predict and calculate loss
5. Compute gradients using `jax.grad`
6. Update weights using simple gradient descent

```
grads = grad(loss_fn)(params, x, y)
params = {k: params[k] - lr * grads[k] for k in params}
```

1. Every 100 epochs: Print accuracy on test data
-

Accuracy Output

Example:

```
Epoch 0, Accuracy: 0.1120
Epoch 100, Accuracy: 0.7945
Epoch 200, Accuracy: 0.8683
...
```

Requirements

```
pip install jax jaxlib tensorflow-datasets
```

File Overview

- `smalljaxproject.py` : Main script containing model, training, and evaluation
- `README.md` : Documentation (this file)

Author

Tejas S

What to Try Next?

- Add batching instead of one sample per step
- Add more hidden layers
- Use `jax.jit` to speed up training
- Visualize weights and predictions
- Build a Gradio web interface